

# cleaning

May 1, 2025

## 1 Bitcoin Price Predictor – AI306

```
[1]: !pip install kagglehub
```

```
Requirement already satisfied: kagglehub in  
/home/mohammed/anaconda3/envs/Crypto/lib/python3.12/site-packages (0.3.4)  
Requirement already satisfied: packaging in  
/home/mohammed/anaconda3/envs/Crypto/lib/python3.12/site-packages (from  
kagglehub) (24.1)  
Requirement already satisfied: requests in  
/home/mohammed/anaconda3/envs/Crypto/lib/python3.12/site-packages (from  
kagglehub) (2.32.3)  
Requirement already satisfied: tqdm in  
/home/mohammed/anaconda3/envs/Crypto/lib/python3.12/site-packages (from  
kagglehub) (4.67.1)  
Requirement already satisfied: charset-normalizer<4,>=2 in  
/home/mohammed/anaconda3/envs/Crypto/lib/python3.12/site-packages (from  
requests->kagglehub) (3.3.2)  
Requirement already satisfied: idna<4,>=2.5 in  
/home/mohammed/anaconda3/envs/Crypto/lib/python3.12/site-packages (from  
requests->kagglehub) (3.7)  
Requirement already satisfied: urllib3<3,>=1.21.1 in  
/home/mohammed/anaconda3/envs/Crypto/lib/python3.12/site-packages (from  
requests->kagglehub) (2.2.3)  
Requirement already satisfied: certifi>=2017.4.17 in  
/home/mohammed/anaconda3/envs/Crypto/lib/python3.12/site-packages (from  
requests->kagglehub) (2024.8.30)
```

```
[2]: import kagglehub  
import shutil  
import os  
  
import numpy as np  
import pandas as pd
```

## 1.1 Data Preprocessing

### 1.1.1 Downloading the dataset

```
[3]: # Download latest version
path = kagglehub.dataset_download("mczielinski/bitcoin-historical-data")

print("Path to dataset files:", path)
```

Warning: Looks like you're using an outdated `kagglehub` version, please consider updating (latest version: 0.3.12)  
Downloading from  
[https://www.kaggle.com/api/v1/datasets/download/mczielinski/bitcoin-historical-data?dataset\\_version\\_number=222...](https://www.kaggle.com/api/v1/datasets/download/mczielinski/bitcoin-historical-data?dataset_version_number=222...)

100%| | 113M/113M [00:12<00:00, 9.65MB/s]

Extracting files...

Path to dataset files:  
/home/mohammed/.cache/kagglehub/datasets/mczielinski/bitcoin-historical-data/versions/222

### 1.1.2 Copying the dataset to the project directory

```
[4]: # Source and destination paths
source_path = os.path.join(path, 'btcusd_1-min_data.csv')
destination_path = './data/btcusd_dataset.csv'

# Create destination directory if it doesn't exist
os.makedirs('./data', exist_ok=True)

# Only copy if the file doesn't already exist
if not os.path.exists(destination_path):
    shutil.copy(source_path, destination_path)
    print("File copied successfully.")
else:
    print("File already exists. Skipping copy.")
```

File already exists. Skipping copy.

### 1.1.3 Reading dataset

```
[5]: data = pd.read_csv(destination_path)
```

/tmp/ipykernel\_3253/1927658998.py:1: DtypeWarning: Columns (6) have mixed types.  
Specify dtype option on import or set low\_memory=False.

```
data = pd.read_csv(destination_path)
```

```
[6]: data.head()
```

```
[6]:      Timestamp  Open  High  Low  Close  Volume      datetime
0  1.325412e+09  4.58  4.58  4.58  4.58      0.0  2012-01-01 10:01:00+00:00
1  1.325412e+09  4.58  4.58  4.58  4.58      0.0  2012-01-01 10:02:00+00:00
2  1.325412e+09  4.58  4.58  4.58  4.58      0.0  2012-01-01 10:03:00+00:00
3  1.325412e+09  4.58  4.58  4.58  4.58      0.0  2012-01-01 10:04:00+00:00
4  1.325412e+09  4.58  4.58  4.58  4.58      0.0  2012-01-01 10:05:00+00:00
```

#### 1.1.4 Info about dataset

```
[7]: print('Shape of the dataset: ', data.shape)
```

Shape of the dataset: (7001004, 7)

```
[8]: print(f"info of the dataset: \n{data.info()}\n")
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7001004 entries, 0 to 7001003
Data columns (total 7 columns):
#   Column      Dtype
---  -
0   Timestamp   float64
1   Open        float64
2   High        float64
3   Low         float64
4   Close       float64
5   Volume      float64
6   datetime    object
dtypes: float64(6), object(1)
memory usage: 373.9+ MB
info of the dataset:
None
```

```
[9]: print(f"describe of the dataset: \n{data.describe()}\n")
```

```
describe of the dataset:
count      Timestamp      Open      High      Low      Close  \
count  7.001004e+06  7.001004e+06  7.001004e+06  7.001004e+06  7.001004e+06
mean    1.535443e+09  1.729476e+04  1.730170e+04  1.728760e+04  1.729476e+04
std      1.212619e+08  2.389940e+04  2.390744e+04  2.389117e+04  2.389938e+04
min      1.325412e+09  3.800000e+00  3.800000e+00  3.800000e+00  3.800000e+00
25%      1.430427e+09  4.239100e+02  4.240000e+02  4.237600e+02  4.239300e+02
50%      1.535442e+09  6.575210e+03  6.578515e+03  6.572320e+03  6.575290e+03
75%      1.640457e+09  2.720000e+04  2.720400e+04  2.719600e+04  2.720000e+04
max      1.745542e+09  1.091110e+05  1.093560e+05  1.087940e+05  1.090360e+05
```

	Volume
count	7.001004e+06
mean	5.308327e+00
std	2.253495e+01
min	0.000000e+00
25%	1.815710e-02
50%	4.703309e-01
75%	3.039586e+00
max	5.853852e+03

```
[10]: print(f"null values of the dataset: \n{data.isnull().sum()}\n")
```

```

null values of the dataset:
Timestamp      0
Open           0
High           0
Low            0
Close          0
Volume         0
datetime      218724
dtype: int64

```

```
[11]: print(f"duplicated values of the dataset: \n{data.duplicated().sum()}\n")
```

```

duplicated values of the dataset:
0

```

### 1.1.5 Cleaning Dataset

```

[12]: # Drop the 'datetime' column
data = data.drop(columns=['datetime'])

# view the updated DataFrame
print(data.head())
print(f"\n\nnull values of the dataset: \n{data.isnull().sum()}\n")

```

	Timestamp	Open	High	Low	Close	Volume
0	1.325412e+09	4.58	4.58	4.58	4.58	0.0
1	1.325412e+09	4.58	4.58	4.58	4.58	0.0
2	1.325412e+09	4.58	4.58	4.58	4.58	0.0
3	1.325412e+09	4.58	4.58	4.58	4.58	0.0
4	1.325412e+09	4.58	4.58	4.58	4.58	0.0

```

null values of the dataset:

```

```

Timestamp    0
Open         0
High         0
Low          0
Close        0
Volume       0
dtype: int64

```

```

[13]: # Convert Unix timestamp (seconds since 00:00:00 UTC January 1, 1970) to
      ↪datetime
data['datetime'] = pd.to_datetime(data['Timestamp'], unit='s')

print(data.head())

```

	Timestamp	Open	High	Low	Close	Volume	datetime
0	1.325412e+09	4.58	4.58	4.58	4.58	0.0	2012-01-01 10:01:00
1	1.325412e+09	4.58	4.58	4.58	4.58	0.0	2012-01-01 10:02:00
2	1.325412e+09	4.58	4.58	4.58	4.58	0.0	2012-01-01 10:03:00
3	1.325412e+09	4.58	4.58	4.58	4.58	0.0	2012-01-01 10:04:00
4	1.325412e+09	4.58	4.58	4.58	4.58	0.0	2012-01-01 10:05:00

```

[14]: # ensure the data is continuous and there are no missing values or rows,
      # Reindexes the data to have a row for every minute - even if that minute was
      ↪missing in the original data.
continuous_data = data.set_index('datetime').asfreq('min')
print('data Null/NA Values before fill:', continuous_data.isnull().values.sum())

# fill in and interpolate missing values after re-indexing is done
continuous_data.interpolate(method='time', inplace=True) # Time-based
      ↪interpolation
continuous_data.ffill(inplace=True) # forwards fill missing values

continuous_data.reset_index(inplace=True) # Moves 'datetime' back from the
      ↪index to a regular column
print('data Null/NA Values after fill:', continuous_data.isnull().values.sum())

data = continuous_data.copy()

```

```

data Null/NA Values before fill: 6960
data Null/NA Values after fill: 0

```

```

[15]: first_nonzero_row = data[data['Volume'] > 0].head(1)
print(first_nonzero_row)

```

	datetime	Timestamp	Open	High	Low	Close	Volume
627	2012-01-01 20:28:00	1.325450e+09	4.84	4.84	4.84	4.84	10.0

```
[16]: # Save cleaned data to csv file
data.to_csv('./data/cleaned_data.csv', index=False)
```

### 1.1.6 Data Reduction

```
[17]: data = data.drop(columns=['Volume'])
data.head()
```

```
[17]:
```

		datetime	Timestamp	Open	High	Low	Close
0	2012-01-01	10:01:00	1.325412e+09	4.58	4.58	4.58	4.58
1	2012-01-01	10:02:00	1.325412e+09	4.58	4.58	4.58	4.58
2	2012-01-01	10:03:00	1.325412e+09	4.58	4.58	4.58	4.58
3	2012-01-01	10:04:00	1.325412e+09	4.58	4.58	4.58	4.58
4	2012-01-01	10:05:00	1.325412e+09	4.58	4.58	4.58	4.58

```
[ ]: def create_resampled_dataframe(data, resample_rule='h'):
    """
    Create a resampled DataFrame for a given time resolution.

    Args:
        data (pd.DataFrame): Input DataFrame
        resample_rule (str): Resampling rule (e.g., 'h', 'D', 'W-MON', 'M')

    Returns:
        pd.DataFrame: Resampled DataFrame

    """
    datetime_column = data['datetime']

    if resample_rule == '1min' or resample_rule is None:
        return data

    df = pd.DataFrame()
    df['Timestamp'] = data.set_index(datetime_column)['Timestamp'].
    ↪resample(resample_rule).first()
    df['Open'] = data.set_index(datetime_column)['Open'].
    ↪resample(resample_rule).first()
    df['High'] = data.set_index(datetime_column)['High'].
    ↪resample(resample_rule).max()
    df['Low'] = data.set_index(datetime_column)['Low'].resample(resample_rule).
    ↪min()
    df['Close'] = data.set_index(datetime_column)['Close'].
    ↪resample(resample_rule).last()
    if 'Volume' in data.columns:
        df['Volume'] = data.set_index(datetime_column)['Volume'].
        ↪resample(resample_rule).sum()
```

```

print('Null/NA Values in resample dataframe:', df.isnull().values.sum())
df = df.dropna()
print('Shape of the dataset: ', df.shape)

return df

```

```

[29]: reduced_data = create_resampled_dataframe(data, resample_rule='h')
      reduced_data.head()

```

Null/NA Values in resample dataframe: 0

Shape of the dataset: (116703, 5)

```

[29]:
      Timestamp  Open  High  Low  Close
datetime
2012-01-01 10:00:00  1.325412e+09  4.58  4.58  4.58  4.58
2012-01-01 11:00:00  1.325416e+09  4.58  4.58  4.58  4.58
2012-01-01 12:00:00  1.325419e+09  4.58  4.58  4.58  4.58
2012-01-01 13:00:00  1.325423e+09  4.58  4.58  4.58  4.58
2012-01-01 14:00:00  1.325426e+09  4.58  4.58  4.58  4.58

```

### 1.1.7 Data Transformation (Feature Engineering)

```

[26]: def add_indicators(data):
      """
      Add technical indicators to the DataFrame.

      Args:
          data (pd.DataFrame): Input DataFrame

      Returns:
          pd.DataFrame: DataFrame with added indicators
      """
      # Calculate moving averages to add to the dataframe
      # moving averages are used to capture the trend of the data by averaging
      ↪ the past values over a period
      sma_200 = data['Close'].rolling(window=200).mean()

      # Calculate Average True Range (ATR)
      # ATR is a volatility indicator that measures the true range over a period
      # It works by comparing the highest and lowest prices over a period
      # Calculate True Range (TR)
      high_low = data['High'] - data['Low']
      high_close_prev = abs(data['High'] - data['Close'].shift(1))
      low_close_prev = abs(data['Low'] - data['Close'].shift(1))

```

```

    true_range = pd.concat([high_low, high_close_prev, low_close_prev], axis=1).
    ↪max(axis=1)

    # Calculate ATR using a rolling average of the True Range
    atr_period = 168 # Or any period of choice
    atr_168 = true_range.rolling(window=atr_period).mean()

    data['SMA_200'] = sma_200
    data['ATR_168'] = atr_168

    print('Null/NA Values in engineered dataframe:', data.isnull().values.sum())
    data = data.dropna()
    print('Shape of the dataset: ', data.shape)

    return data

```

```

[ ]: # this is the cleaned, reduced (minute > hourly) and engineered data
engineered_data = add_indicators(reduced_data)
engineered_data.head()

```

Null/NA Values in engineered dataframe: 366

Shape of the dataset: (116504, 7)

```

[ ]:

```

	datetime	Timestamp	Open	High	Low	Close	SMA_200	ATR_168
	2012-01-09 17:00:00	1.326128e+09	6.9	6.9	6.5	6.5	5.83495	0.032560
	2012-01-09 18:00:00	1.326132e+09	6.5	6.6	6.5	6.5	5.84455	0.033155
	2012-01-09 19:00:00	1.326136e+09	6.5	6.6	6.5	6.6	5.85465	0.033750
	2012-01-09 20:00:00	1.326139e+09	6.6	6.6	6.6	6.6	5.86475	0.033750
	2012-01-09 21:00:00	1.326143e+09	6.6	6.6	6.6	6.6	5.87485	0.033750

```

[30]: data = engineered_data.reset_index() # reset the index so that we can use it as
    ↪a column
data.head()

```

```

[30]:

```

	datetime	Timestamp	Open	High	Low	Close	SMA_200	ATR_168
0	2012-01-09 17:00:00	1.326128e+09	6.9	6.9	6.5	6.5	5.83495	0.032560
1	2012-01-09 18:00:00	1.326132e+09	6.5	6.6	6.5	6.5	5.84455	0.033155
2	2012-01-09 19:00:00	1.326136e+09	6.5	6.6	6.5	6.6	5.85465	0.033750
3	2012-01-09 20:00:00	1.326139e+09	6.6	6.6	6.6	6.6	5.86475	0.033750
4	2012-01-09 21:00:00	1.326143e+09	6.6	6.6	6.6	6.6	5.87485	0.033750

```

[31]: # Save cleaned data to csv file
data.to_csv('./data/preprocessed_hourly_data.csv', index=False)

```